

# IN3200/IN4200: Chapter 3

## Data access optimization

### (Part 3)

Textbook: Hager & Wellein, *Introduction to High Performance Computing for Scientists and Engineers*

## Goals of Chapter 3:

- be able to reason about data access and performance using machine/code balance analysis
- be able to recognize and apply data access optimizations
  - loop unrolling
  - loop unroll-and-jam
  - loop blocking
  - loop fusion
  - re-organization of data structure

Three cases of code balance analysis and data access optimization:

- Dense matrix-vector multiply (repetition)
- Matrix transpose
- Sparse matrix-vector multiply

# Mathematical definition of matrix-vector multiply

Square matrix A:  $N$  rows and  $N$  columns of numerical values

Vector B:  $N$  numerical values

Vector C:  $N$  numerical values

Mathematical definition of matrix-vector multiply:  $C = C + A * B$   
such that each value in vector C is calculated as

$$C_i = C_i + \sum_{0 \leq j < N} A_{i,j} * B_j \quad 0 \leq i < N$$

# Dense matrix-vector multiply (repetition)

Here, we consider the case of  $A$  being a “dense” matrix: all its  $N \times N$  numerical values are nonzero.

Storage on a computer:

- Dense square matrix  $A$  as a 2D array,  $N$  rows and  $N$  columns, row-major storage (in C language)
- Vectors  $B$  and  $C$  each as a 1D array of length  $N$

Each value  $A_{i,j}$  is accessed as  $A[i][j]$

## Straightforward implementation & balance analysis


```
for (i=0; i<N; i++) {  
    double tmp = C[i];  
    for (j=0; j<N; j++)  
        tmp = tmp + A[i][j]*B[j];  
    C[i] = tmp;  
}
```

- Total number of floating-point (FP) operations:  $2N^2$
- Memory traffic:  $N^2$  loads for 2D array A,  $N$  loads &  $N$  stores for 1D array C
- How many loads are associated with 1D array B?
  - Small cache  $\rightarrow$  array B is loaded  $N$  times  $\rightarrow N^2$  memory loads
  - Large cache  $\rightarrow$  array B is loaded only once  $\rightarrow N$  memory loads

Code balance for the small-cache case:

$$\frac{N^2 + N^2 + 2N}{2N^2} = 1 + \frac{1}{N}$$

# Illustration of array B being loaded $N$ times



**Figure 3.11:** Unoptimized  $N \times N$  dense matrix vector multiply. The RHS vector is loaded  $N$  times.

# How to reduce memory traffic for small-cache case?

## *m*-way unroll-and-jam:


- Unroll the outer loop *m* times
- Fuse the *m* inner loops

```
for (i=0; i<N; i+=m) {  
    for (j=0; j<N; j++) {  
        C[i+0] += A[i+0][j]*B[j];  
        C[i+1] += A[i+1][j]*B[j];  
        // ...  
        C[i+m-1] += A[i+m-1][j]*B[j];  
    }  
}  
// remainder code in case (N%m)>0 ....
```

- *m*-fold reuse of each  $B[j]$  from register
- Number of memory loads for array B:  $N^2/m$  (for small-cache case)
- Size of *m* shouldn't be too large, to avoid too high *register pressure*



# Illustration of the effect of unrolling



**Figure 3.12:** Two-way unrolled dense matrix vector multiply. The data traffic caused by reloading the RHS vector is reduced by roughly a factor of two. The remainder loop is only a single (outer) iteration in this example.

For the small-cache case, unroll-and-jam will result in the following improved code balance:

$$\frac{N^2 + \frac{N^2}{m} + 2N}{2N^2} = \frac{1}{2} + \frac{1}{2m} + \frac{1}{N}$$

# Matrix transpose

The *transpose* of an  $N$ -by- $N$  matrix  $B$  is another  $N$ -by- $N$  matrix  $A = B^T$  such that  $A_{i,j} = B_{j,i}$ .

```
for (j=0; j<N; j++)  
    for (i=0; i<N; i++)  
        A[j][i] = B[i][j];
```

It is assumed that  $A$  and  $B$  are 2D arrays in row-major storage.

(Note: The matrix transpose example in the textbook (Section 3.4) is programmed in Fortran and assumes column-major storage!)

In the above code, values are loaded from  $B$  in the order

$$B[i][j], \quad B[i+1][j], \quad B[i+2][j], \quad \dots$$

These large jumps in memory can lead to poor cache line utilization.


# Loop unrolling applied to matrix transpose

## *m*-way unroll-and-jam:

- Unroll the outer loop *m* times
- Fuse the *m* inner loops

```
for (j=0; j<N; j+=m) {  
    for (i=0; i<N; i++) {  
        A[j+0][i] = B[i][j+0];  
        A[j+1][i] = B[i][j+1];  
        // ...  
        A[j+m-1][i] = B[i][j+m-1];  
    }  
}
```

# Illustration of unrolled matrix transpose



**Figure 3.13:** Two-way unrolled “flipped” matrix transpose (i.e., with strided load in the original version).

B

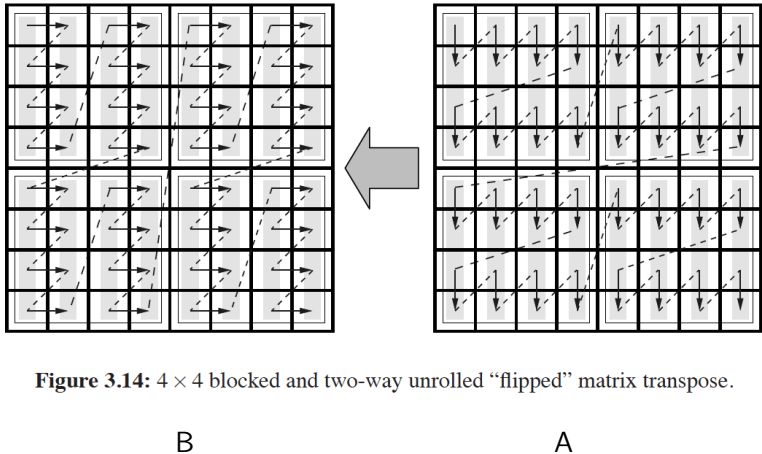
A

## Loop blocking + unrolling applied to matrix transpose

```
for (jj=0; jj<N; jj+=b) {  
    jstart = jj; jstop = jj+b-1;  
    for (ii=0; ii<N; ii+=b) {  
        istart = ii; istop = ii+b-1;  
  
        for (j=jstart; j<=jstop; j+=m) {  
            for (i=istart; i<=istop; i++) {  
                A[j+0][i] = B[i][j+0];  
                A[j+1][i] = B[i][j+1];  
                // ...  
                A[j+m-1][i] = B[i][j+m-1];  
            }  
        }  
    }  
}
```

Blocking improves locality for accessing B.

# Loop blocking + unrolling applied to matrix transpose




**Figure 3.14:**  $4 \times 4$  blocked and two-way unrolled “flipped” matrix transpose.

When most of the numerical values of matrix  $A$  are zero, it is called a *sparse* matrix.

- It will be a waste of float-point operations if we still use the straightforward (dense matrix-vector multiply) implementation
- It will also be a waste of storage if we store a sparse matrix as a 2D array



# Illustration of sparse matrix-vector multiply



**Figure 3.15:** Sparse matrix-vector multiply. Dark elements visualize entries involved in updating a single LHS element. Unless the sparse matrix rows have no gaps between the first and last nonzero elements, some indirect addressing of the RHS vector is inevitable.

# Basic idea for saving storage and computation

- Store only the nonzero values of  $A$ 
  - 2D-array format can no longer be used
  - many sparse storage formats are possible
  - all sparse formats must somehow store the row  $i$  and column  $j$  of every *nonzero* value  $A_{i,j}$
- Avoid multiplications with zero
  - If  $N_{\text{nz}} (\ll N^2)$  denotes the number of nonzero values in a sparse matrix  $A$ , then we only need  $2N_{\text{nz}}$  floating-point operations (instead of  $2N^2$  FP) for a sparse matrix-vector multiply

# Coordinate storage (COO) format

Three 1D arrays of length  $N_{\text{nz}}$ :

- `val`, stores all the nonzero values of the sparse matrix
- `row_idx`, stores the row positions of the nonzero values
- `col_idx`, stores the column positions of the nonzero values

```
for (int k=0; k<Nnz; k++)
```

```
    C[row_idx[k]] = C[row_idx[k]] + val[k]*B[col_idx[k]];
```

- Single loop over all nonzeros (of length  $N_{\text{nz}}$ )
- Accesses to arrays `val`, `row_idx` and `col_idx` are with stride one (good spatial locality)
- Accesses to arrays `B` and `C` are indirect (via `row_idx` and `col_idx`) and can be completely irregular (spatial and temporal locality depends on row and column positions)

# Code balance analysis of matrix-vector multiply with COO

```
for (int k=0; k<Nnz; k++)  
    C[row_idx[k]] = C[row_idx[k]] + val[k]*B[col_idx[k]];
```

Assume that each entry in `row_idx` and `col_idx` is half a word.

**Best-case scenario:** entire `B` and `C` arrays are cached, needing in total only  $2N$  loads and  $N$  stores:

$$\frac{N_{\text{nz}}(1 + 0.5 + 0.5) + N + 2N}{2N_{\text{nz}}} = 1 + \frac{3}{2} \frac{N}{N_{\text{nz}}}$$

**Worst-case scenario:** `B[col_idx[k]]` and `C[row_idx[k]]` need to be loaded from and stored to memory every single time, and only one value is used per cacheline:

$$\frac{N_{\text{nz}}(1 + 0.5 + 0.5) + 3N_{\text{nz}} \frac{\text{cacheline size}}{\text{word size}}}{2N_{\text{nz}}} = 1 + \frac{3}{2} \frac{\text{cacheline size}}{\text{word size}}$$

# Compressed row storage (CRS) format

Idea: *grouping nonzeros by rows* and thus fewer accesses to array C




Figure 3.16: CRS sparse matrix storage format.

Note: In the above figure from Chapter 3, arrays row\_ptr and col\_idx contain 1-based indices due to Fortran programming.

Three arrays:

- 1D array val, of length  $N_{nz}$ , stores all the nonzero values of the sparse matrix
- 1D array col\_idx, of length  $N_{nz}$ , records the original column positions of all the nonzero values
- 1D array row\_ptr, of length  $N + 1$ , contains the indices at which

# Implementation of matrix-vector multiply using CRS format

```
for (i=0; i<N; i++) {  
    double tmp = C[i];  
    for (j=row_ptr[i]; j<row_ptr[i+1]; j++)  
        tmp = tmp + val[j]*B[col_idx[j]];  
    C[i] = tmp;  
}
```

- There is a long outer loop (of length  $N$ )
- The inner loop can be very short
- Access to array `C` will be well optimized by compiler
- Access to array `val` is with stride one (perfect situation)
- Access to array `B` is indirect (via `col_idx`) and can be irregular

# Code balance analysis of matrix-vector multiply with CRS

Assume that each entry in `row_ptr` and `col_idx` is half a word.

**Best-case scenario:** entire `B` array is cached, needing only  $N$  loads:

$$B_c = \frac{N_{\text{nz}}(1 + 0.5) + 0.5N + N + 2N}{2N_{\text{nz}}} = \frac{3}{4} + \frac{7}{4} \frac{N}{N_{\text{nz}}}$$

Note: In Chapter 3.6.1, page 87, the estimated code balance  $B_c = \frac{5}{4}$  is not optimal.

**Exercise:** What is the code balance for the *worst-case scenario*, where `B[col_idx[j]]` must be loaded from memory every single time, and only one value is used per cacheline?

- Continue using CRS format, but with suitable permutations (to reduce the actual memory traffic associated with array B)
- Use the JDS format (which targets vector processors) with further optimization (see Sections 3.6.1 & 3.6.2)